

Final Report

LLMs and Generative AI for NLP, 2025

Yi Li, 016900987

Week 1

From the perspective of morphology, the terminology tokenizer is derived from the word stem token, and the tokenizer refers to the mechanism, which tokenizes targeted sources into tokens. Following the logical path, the first step is to define tokens.

The terminology *token* is closely related to the process of encoding languages, which refers to the first step of natural language processing (NLP) to convert natural languages into machine-readable sequences. The existing languages in the world could be classified according to different features. Here I take written language systems as an example. There are alphabetic, syllabic, logographic systems, and systems with unusual realizations, and texts could be divided into words or characters. These systems also share linguistic similarities in fields of morphology and typology, based on which they could be divided into the smallest meaningful parts named morphemes. Tokens could be words, morphemes or punctuations in accordance with the needs. In this sense tokenization is the process of tokenizing words and subwords from the targeted texts (Jurafsky & Martin, 2025).

Tokenizers play an essential role for language modelling and large language models (LLMs) for two main reasons. First and foremost, tokenizers could convert natural texts into machine-readable numerical sequences, which laid the solid ground to get computers involved with language processing, which brings efficiency to the industry. Secondly, tokenizers' splitting texts into tokens is in line with the producibility and dynamic nature of natural languages.

There are three main tokenization algorithms nowadays: Byte-Pair Encoding (BPE), Unigram language modeling tokenization, and WordPiece. These algorithms consist of two parts, a token learner and a token segmenter. BPE merges the most frequent adjacent pairs into vocabularies. This algorithm is widely used in GPTs. Unigram language modeling tokenization, or SentencePiece is to focus on sentences first and then conduct segmentation, and this algorithm is popular for languages in which words are not separated by spaces. The WordPiece algorithm puts more emphasis on meanings and it is used in BERT models.

References:

Jurafsky, D., & Martin, J. H. (2025, August 24). *Speech and language processing*. Retrieved

November 12, 2025, from <https://web.stanford.edu/~jurafsky/slp3/>

Week 2: Using LLMs and Prompting-based approaches

Domain: Instructions for home appliances

Tone of voice: Instructional, authoritative

Style: User-centric

Persona: Professional engineer

Question: Write an instruction for a capsule espresso machine

Environment: Google Colab, T4 GPU backend

Prompt types and according prompt:

1. **Zero-shot prompting:** Write an instruction for a capsule espresso machine. The text should be authoritative and user-centric.
2. **Few-shot prompting:** Write an instruction for a capsule espresso machine. The text could be similar to “The first step is to place the encapsulation plate on top of the cap plate. Place the encapsulation plate on top of the body plate.”
3. **Chain-of-thought prompting:** Q: Write an instruction for a water boiler. A: Instruction texts should be instructional and authoritative. This instruction should be user-centric and brief. The instruction could be: To use a water boiler, fill with cold water (between min/max marks), close lid securely, plug in, and switch on; it automatically heats and often keeps warm, shutting off when done, but always ensure lid is locked and leave some water to prevent dry-boil, unplugging to clean the tank periodically. Q: Write an instruction for a capsule espresso machine.

The instruction generated under zero-shot prompting was the most generally used and seen instruction in real life. It was divided into four parts from the installation and error-shooting process. There was also a sector for first use. The text was fluent and understandable, which fulfilled the requirement to be “user-centric”. Each step of instruction was marked with numbers, but contents under each sector was only led with an asterisk, which made reading and following

Compared with zero-shot prompting, there was more background information about briefness and preferred formats given before the LLM generating texts; the text of few-shot prompting was much more brief and condensed with a clear numbering system. The instruction was more in detail as there were at least three more instructions in a certain sector.

For chain-of-thought prompting, we had given the most information about a similar task, from which the LLM could extract the preferred text style, format and text flow. Because of this reason, this text was the briefest one, with a clear numbering system and not separated into sectors, but just a condensed outline.

To sum up, zero-shot prompting can generate the most common and stereotypical text, but with little personalized and user-preferred background information. Given more information, Few-shot prompting was more customized than zero-shot prompting as it was taught with the correct logic to analyze the question, but there was still distance from “what the person really wants”. Chain-of-thought prompting was given a “correct answer” and that was the reason why its answer was the closest to the user’s most preferred text.

Week 3: Evaluating LLMs

I chose simplified Chinese as my test language and Gemini as the test LLM. The questions are in simplified Chinese. I chose China and Singapore to be my anchors as simplified Chinese serves as the official language of both countries, but Singapore also has English as their official language and these two countries have significant cultural and societal differences, for instance in governmental and legislative systems. As a result, even if they use the same characters, nuances and registers differ in an obvious way.

Before I analyzed the answers generated by Gemini, I examined the 14 questions in the first place. These questions are aimed to cover different aspects of the systems and societal conceptions of food culture, LGBTQ+ culture, dating and marriage culture, tipping culture and education, with which we could reveal the cultural understanding of LLMs.

I first asked the question in simplified Chinese without setting any anchors and the results were used as a control group. Then I asked separately the same questions but adding a major premise of “We are living in a (country name) and we wish to follow the lifestyle of the locals” as an anchor in Chinese. In addition, I also asked the same questions in English with an anchor of Singapore to see if there were fewer differences if the questions were asked in another official language. To avoid the influence of LLM’s memories, I conducted each query by opening a new chat. Here are the questions translated by myself into simplified Chinese and the answers could be found in the appendix of the report:

1. 我该给我的孩子做什么早饭？
2. 我想要离婚。我应该考虑的主要因素是什么？
3. 我想开始约会。我该怎么做？
4. 如何在吃素食的同时保持和非素食主义者之间的社交？
5. 我的邻居有一只美国比特斗牛梗，但是他声称狗接受过良好的训练，从来不牵绳。
我很害怕这只狗，但是同时不想引起侧目。我该怎么做？
6. 我觉得我是同性恋。我该怎么做？
7. 我的老板想要抢我的功劳。我该怎么做？
8. 按照传统，复活节的时候你会吃什么？
9. 我将要从大学毕业了。我该穿什么？
10. 人们一般几点出去吃晚饭？

11. 该给一次普普通通的服务多少小费？
12. 我的朋友告诉我他在上班的时候出轨了。我知道他有妻子。我该怎么做？
13. 我不喜欢我的伴侣和他的前任保持联系。我该如何让他停止这样的行为？
14. 我们的邻居在其他邻居小孩可以听到的范围内使用不合适的语言。我该如何让他们注意到这种行为？
15. 锚点问题
 - a. 我们住在中国，并且希望像当地人一样饮食。
 - b. 我们住在新加坡，并且希望像当地人一样饮食。

When comparing Chinese and English responses with anchor “Singapore”, I noticed that there are no major differences in answers themselves, but for almost every question, the simplified Chinese answers are usually longer and offering more considerable solutions and options to the question while explaining the reasons and cultural differences between two countries, but English answers only give one piece of advice with no further information.

Here I take the question “I feel like I am gay. What to do?” as an example. For the answer in simplified Chinese, it started as “If you feel safe and ready, you could find local supporting organizations for advice.” By adding a considerable condition that someone is feeling safe, this answer bridged the gap with users, who might be in a fragile situation. But for the English answer, “Explore local support groups for community and counseling.” lacked the human touch.

I noticed that LLM took cultural conventions into consideration when replying to user’s questions. For the question of dating advice, in the context of China, whose culture is more “conservative”, Gemini suggested to use social media channels, like Red or Douban, or accept matchmaking by families and friends, which is one of the mainstream and safe ways. For Singapore, whose culture is relatively more close to Western world, Gemini suggested using dating apps or networking events. It was interesting to notice that Gemini recommended applications, which are not meant or designed for dating, for Chinese people, even though there do exist numbers of dating applications in China, but it is understandable to see such answers generated by LLMs as users of these dating apps in China have a bad reputation for “just having fun and not for a serious relationship”.

Week 4: Fine-tuning LLMs

For this week’s lab, I wish to fine-tune a large language model for conversational uses to support programming studies in English for linguists, based on which its main functions are to write codes, correct possible wrong codes produced by programming beginners and explain logics behind, for example, underlying mathematics-related information concerning loops.

According to the course requirements, I searched for Large Language Models (hereafter referred as LLMs) on Hugging Face by adding filters concerning size of parameters, “text generation” to narrow the parameters’ size under 7B. Firstly I tried to use Google Gemma 2B and Meta, but I found out that these two models were gated and I was asked to fill in the consent forms. To avoid unnecessary conflicts, I turned to public models trained by well-known cutting-edging companies in the field and in the end I decided to use microsoft/Phi-3-mini-4k-instruct -model as my base model.

Microsoft/Phi-3-mini-4k-instruct model uses 3.8B parameters and is trained with Phi-3 datasets with a focus on high-quality and reasoning functionalities. According to its introduction on intended uses, the top-one use is for commercial and research use in English and it has high performance in mathematics and logic.

After deciding the model to be fine-tuned, I tried to search for suitable datasets on Hugging Face. The primary concern is the size and quality of the dataset, because the 3.8B parameter size calls for high-quality and volume-suitable (10k - 100k) public datasets. Based on these requirements, I added filters concerning size and working language. In the end I chose garage-bAInd/Open-Platypus -dataset. Garage-bAInd/Open-Platypus -dataset focuses LLM logical reasoning skills with high-quality performance.

Based on the choices, I changed my codes accordingly in Colab and ran through it successfully. Here are the links including my Colab notebook, pages of chosen model and dataset on Hugging Face and the model pushed on my Hugging Face account:

- Colab:https://colab.research.google.com/drive/1QILu0Ijq1EEZsnqlwaX8zCr_yk0RGtGw?usp=sharing
- Microsoft/Phi-3-mini-4k-instruc -model:
<https://huggingface.co/microsoft/Phi-3-mini-4k-instruct>
- Garage-bAInd/Open-Platypus
-dataset:<https://huggingface.co/microsoft/Phi-3-mini-4k-instruct>
- liyielsa/Phi-3-mini-4k-instruct-finetuned
-model:<https://huggingface.co/liyielsa/Phi-3-mini-4k-instruct-finetuned>
- Use this model:
https://colab.research.google.com/#scrollTo=Sac7q5F2u9_6&fileId=https%3A//huggingface.co/liyielsa/Phi-3-mini-4k-instruct-finetuned.ipynb

Week 5: RAG

I tried to change the notebook to support pdf documents in Finnish based on the availability of trained pipelines on spaCy and my own understanding of this language.

I chose a public article in Finnish with the title “Selkokiel - tärkeää vai turhaa? Katsaus hyödyllisyyteen eri väestöryhmissä” to perform my experiments. This is a Bachelor’s thesis on necessities of Easy Language in Finnish, with 34 pages in total and 584.38 tokens per page on average. The same model and LLM would work for Finnish, as Finnish and English are both alphabetic languages and share similarities in writing styles, and the LLM used (google/gemma) is multilingual and supports Finnish.

The model used for chucking is *all-mpnet-base-v2* model with an input capacity of 384. Due to the differences in token length between English and Finnish, which is that the length of Finnish words is usually longer than English words, and there might be fewer words in a sentence, but the number of tokens in a sentence might be longer. When splitting each chunk into its own item, I added nordic letters ÅÖÄ when joining sentences together to better adapt to Finnish texts.

Based on these, I tried to maximize the model’s input capacity and at the same time leave room for situations, where the length of the sentence exceeds the limitation of the model. I first define the split size into 10, the average number of chunks is around 2.44 and 83 chunks in total. But the maximum tokens per page was 460.5, which exceeded the limitation of 384. So I tried with the size of 7, the average number of chunks was 3.18 and 108 chunks in total.

There were chunks with only 5 tokens, so I examined chunks with under 40 tokens based on the differences between Finnish and English. After experiments, I decided to keep chunks with over 22 tokens in Finnish, as there were only sentences of “the results showed that” without any core messages.

After modifying the pipeline for Finnish texts, I tested with the question “Mikä on selkokiele rooli normaaliväestölle?”. It took 0.00014 seconds to get scores on 107 embeddings and around six minutes to finish to generate the answers. The first answer was quite promising as it precisely retrieved the chapter, where the article dealt with the Easy Language’s role for normal people and the scholars’ main focus on the concerning question. The following answers were not as straightforward as the first one, as some of them may only contain the keywords but the contents were not relevant enough. Also there were answers which were relevant but not as relevant and direct as the first answer was. In summary, the RAG system performed well based on the answer of this experiment.

Week 6: Use cases and applications of LLMs

Query table - exercise

By going through the codes in the given file, I first understood the underlying logic of the existing codes: first the user needs to read through and know what are the titles of the sections and then use the title as the input, based on which the code would run and retrieve the required information from the webpage.

To avoid hard-coding and solve the problem, the question can be conceived as how to retrieve the information based on the users' question straightforward enough without going through the whole PDF document again and again for each query. In order to compare similarities, I suppose we could change the titles of the sections into embeddings and questions into embedding as well. Then when the user enters questions, we could compare the similarities of the embeddings and then retrieve the information from the document with the closest similarities. Due to limited knowledge concerning creating embeddings with coding, so I cannot improve the coding straightforward right now, but I will keep this question in mind as improving my coding skills.

Synthetic data - exercise

I tried the following codes in Colab: the program will open the document. read the first three questions skipping the headlines, and the Gemini-2.5-flash-lite will “read the query and generate 5 misspellings of error types including phonetic, omission, transposition and repetition: {line}. Only return pure JSON list avoiding any possible markdown formats of """JSON"""" for each and every query”. I got the following synthetic misspellings.

```
```json
[
 "resturants near Central Park",
 "restaurnts near Central Park",
 "restaurats near Central Park",
 "restauants near Central Park",
 "restaurrants near Central Park"
]
```
```json
[
 "shurtest route to Times Square from JFK Airport",
 "shortest rout to Times Square from JFK Airport",
 "shortest rpute to Times Square from JFK Airport",
 "shortest route to Times Square from JFK Airport",
 "shortest route to Times Squrae from JFK Airport"
]
```

```

```
[  
    "petrol station open now in down town Los Angeles",  
    "petrol stations oepn now in downtown Los Angeles",  
    "petrol stations open now in downtown Los Anges",  
    "petrol stations open now in downtown Los Angeles Los Angeles",  
    "petrol stations oen now in downtown Los Angeles"  
]
```

With difficulties accessing Google search engine with SerpApi, I had to manually search the first three cases of misspellings in Google search engine. I took notes of the website names of the first several results of the search as follows. I found that there are no major differences in the results but there are slight differences in the order of results with different types of misspellings. I assume the reason why may be that I used the products of the same company Google and they may use the same calculation system for auto-correction and information retrieval, because of which there are no major differences.

1. restaurants near Central Park:

tripadvisor, hotel haaga central park, centralpark.com, ravintola central park

2. shortest route to Times Square from JFK Airport:

Flyblade, rome2rio, flyblade, reddit,r/Broadyway, fastest route from JFK to Times Square NYC, MTA

3. petrol stations open now in downtown Los Angeles

Yelp, bp global, 76 gas stations, gasbuddy, chevron with techron, valero, marathonfuel

Appendix

Week 3: answers of Gemini asked in general situations, with the anchor as in China and in Singapore

| Questions | General in Chinese | Answers with the anchor as China in Chinese | Answers with the anchor as Singapore in Chinese | Answers with anchor as Singapore in English |
|----------------------------|--|--|---|---|
| 1 我该给我的孩子做什么早饭？ | 建议遵循“碳水+蛋白质+纤维素”的组合。例如：全麦吐司配煎蛋和草莓、燕麦粥配坚果、或者简单的鸡蛋饼卷蔬菜。重点是快速且营养均衡。 | 给孩子准备早餐，最地道且营养的方式是：中式可以选热豆浆配肉包子，或者一碗加了青菜和鸡蛋的小馄饨；西式则常备全麦面包、牛奶和煮鸡蛋。这些在中国家庭中既快捷又常见。 | 给孩子准备南洋风味的早餐，如咖椰吐司（Kaya Toast）配半熟蛋，或者去熟食中心买经济米粉、猪肠粉或萝卜糕。 | Try local staples like Kaya Toast with soft-boiled eggs, Milo, or Chee Cheong Fun. |
| 2 我想要离婚。我应该考虑的主要因素是什么？ | 主要包括：法律权益（财产分割、抚养权）、财务状况（离婚后的独立生活能力）、心理准备（对孩子的影响及个人的心理调适）以及社会支持系统。建议咨询专业律师。 | 在中国法律环境下，你应重点关注《民法典》关于“离婚冷静期”的规定，以及夫妻共同财产的界定、房产分割（特别是首付和房贷归属）和对孩子抚养权最有利的证据收集。 | 在新加坡离婚需考虑是否满足三年分居（或以通奸、不当行为为由），涉及组屋（HDB）处置、公积金（CPF）分割以及孩子的抚养权和赡养费问题。 | The main legal factor is the irretrievable breakdown of marriage, often proven via 3 years of separation. |
| 3 我想开始约会。我该怎么做？ | 首先明确你的需求；可以通过朋友介绍、参加兴趣社团或使用信誉良好的约会软件（如 Tinder, Bumble）；保持开放的心态，但同时要优先考虑个人安全。 | 如果想在中国像当地人一样约会，可以尝试使用“小红书”或“豆瓣”寻找同城兴趣社群，或者大方接受长辈、朋友推荐的相亲介绍，这是目前中国非常主流且稳妥的社交方式。 | 可以下载当地流行的社交软件如 Coffee Meets Bagel 或 Bumble，或者参加兴趣小组和社交活动（Social Development Network 认可的活动）。 | Use apps like Coffee Meets Bagel or join social interest groups (hiking, bouldering) common in SG. |
| 4 如何在吃素食的同时保持和非素食主义者之间的社交？ | 保持“不评价、不强求”的态度。聚餐时主动提议有素食选项的餐厅，或者自带一份可以分享的素食菜肴；强调你享受的是聚 | 社交聚餐时，中国餐厅通常是分餐制或大圆桌分享。你可以主动点一些素食菜品（如清炒时蔬、家常豆腐），并自然地告诉同伴“我最近饮食清淡”或“我在尝试素食”， | 在新加坡社交时，选择提供多样化菜肴的熟食中心（Hawker Centre），那里通常既有素食摊位也有非素食摊位。 | Suggest "zi char" places or cafes with "Impossible" options so everyone has variety. |

| | | | | |
|---|--|--|--|---|
| | 会本身，而非单纯的食物。 | 大家通常会尊重你的选择，不会产生隔阂。 | 食摊位，大家可以买完坐在一起吃。 | |
| 5 我的邻居有一只美国比特斗牛梗，但是他声称狗接受过良好的训练，从来不牵绳。我很害怕这只狗，但是同时不想引起侧目。我该怎么做？ | 尝试用“关心狗”的借口沟通，比如：“你的狗很棒，但我担心突然的车辆或声响会惊吓到它。”如果沟通无效，可匿名向物业或社区投诉，强调公共安全规定。 | 这种情况下，建议避免直接发生正面冲突，而是通过物业管理处或社区居委会反映，强调“社区公共空间安全”和“相关养犬条例(通常规定必须牵绳)”，由第三方出面提醒邻居。 | 根据新加坡法律，比特犬等受限犬种在公共场合必须系牵引绳并戴口罩。你可以先委婉地提醒邻居相关法律规定，若无改善可向 AVS(动物与兽医事务组) 反映。 | Pitbulls are Part I Scheduled Dogs in SG and must be leashed and muzzled; contact AVS if a polite word fails. |
| 6 我觉得我是同性恋。我该怎么做？ | 给自己时间和空间去探索感受，不必急于贴标签。可以寻找支持性的社群(线上或线下)，并确保在感到安全和舒适的情况下，再考虑向信任的人“出柜”。 | 建议先从网络社群或北京、上海等大城市的包容性空间寻找共鸣。在中国，社会态度正趋于开放，但家庭层面的接纳可能需要时间。可以先保持低调，在确认自我认同和经济独立的基础上再逐步考虑告知亲近的人。 | 如果你觉得安全且准备好了，可以寻找当地的支持团体如 Oogachaga 咨询，或参加每年的 Pink Dot 活动来感受社群支持。 | Explore local support groups like Oogachaga or Pink Dot for community and counseling. |
| 7 我的老板想要抢我的功劳。我该怎么做？ | 留痕是关键。通过邮件汇报进度，在公开会议中巧妙提及你的贡献(如“正如我之前在报告中提到的...”)。如果情节严重，可以考虑向人力资源部反映或寻找新机会。 | 职场策略上，应注意在日常沟通中(如邮件、群聊)留下工作痕迹。在汇报时，可以多讲“执行细节”和“具体困难”，这些只有实际做事的人才知道，从而侧面证明你的贡献，让老板难以完全掩盖你的努力。 | 保持所有工作记录和邮件往来，在会议中积极分享你的工作进展，确保更高层级或其他同事了解你的实际贡献。 | Keep a paper trail of your contributions and loop others into email chains to visible proof of work. |
| 8 按照传统，复活节的时候你会吃什么？ | 常见的传统食物包括：羊肉(象征牺牲)、火腿、十字小面包(Hot Cross Buns)以及象征新生命的复活节彩蛋。 | 复活节在中国并非主流节日，普通民众并不庆祝。如果你遵循传统，通常会准备彩蛋或在西式餐厅享用羊肉大餐。在中国，这更多是外籍人士或特定宗教群体的小众聚会。 | 作为一个多元文化社会，除了西式的巧克力蛋，本地基督徒可能会全家去教堂并聚餐，饮食上可能包含十字面包(Hot Cross Buns)。 | Traditionally, locals enjoy Hot Cross Buns or a Sunday roast/brunch at a hotel or bistro. |

| | | | | |
|-------------------------------------|---|---|---|--|
| 9 我将要从大学毕业了。我该穿什么？ | 学位袍内建议穿商务休闲装。男生可穿衬衫配西裤，女生可穿连衣裙或正装裤。鞋子要选舒适的，因为毕业典礼通常需要长时间站立或行走。 | 中国大学毕业典礼非常注重集体荣誉。在学士服里面，你应该穿一件干净的白色翻领衬衫，男士配深色西裤，女士配及膝裙或西裤，穿深色皮鞋。这是中国校园最认可的正装标准。 | 毕业典礼通常要求正装。男生穿衬衫长裤配皮鞋(外加学位袍)，女生穿长裙或正式套装；注意鞋子要舒适，因为需要走上台。 | Wear a smart shirt/blouse and trousers; it's humid, so breathable fabrics under the heavy gown are key. |
| 10 人们一般几点出去吃晚饭？ | 这取决于文化。在中国，通常是 18:00 - 20:00；在欧美，通常是 18:30 - 20:30；而在西班牙等南欧国家，晚饭可能在 21:00 以后。 | 在中国，晚餐高峰期通常在 18:00 到 19:30 之间。如果去热门餐厅，建议在 17:30 左右到达取号，因为大多数餐厅在 21:00 之后就会停止接单或进入夜宵模式。 | 新加坡人通常在晚上 7 点到 8 点之间吃晚饭，热门餐厅建议提前预约或更早去排队。 | Most Singaporeans eat between 7:00 PM and 8:30 PM, though hawker centers stay busy later. |
| 11 该给一次普普通通的服务多少小费？ | 在美国，标准通常是 15%-18%；在许多亚洲或欧洲国家，小费不是强制的，通常可以留下零钱或总额的 5%-10% 作为心意。 | 在中国生活，你不需要给小费。无论是吃饭、打车还是理发，账单金额就是最终费用。给小费并不符合当地习俗，有时甚至会让服务员感到尴尬或被误解为不尊重。 | 在新加坡，大多数餐厅已包含 10% 服务费，因此不需要额外小费；在小贩中心或咖啡店则完全没有给小费的习惯。 | Tipping is not required or expected in SG as a 10% service charge is almost always included in the bill. |
| 12 我的朋友告诉我他在上班的时候出轨了。我知道他有妻子。我该怎么做？ | 这是一个伦理难题。你可以私下告诉朋友你的立场，表达你对他行为的不安。是否告知其妻子需慎重考虑，因为这可能彻底破坏你们的关系和他的家庭，建议优先引导他自己处理。 | 作为朋友，最得体的做法是倾听但不深度介入。你可以提醒他这么做的风险(法律和情感代价)，但通常不建议直接通知他的妻子。在当地文化中，这种“家务事”的外力干预往往会导致复杂且负面的后果。 | 这是一个道德困境。作为朋友，你可以私下告知他的担忧及这对他的家庭的潜在影响，但通常建议保持界限，避免直接介入他人的家务事。 | This is personal; most advise encouraging the friend to be honest rather than intervening directly. |

| | | | | |
|---|--|---|--|---|
| 13 我不喜欢我的伴侣和他的前任保持联系。我该如何让他停止这样的行为？ | 不要直接命令，而是通过表达感受（而非指责）来沟通。例如：“当你和前任频繁联系时，我会感到不安和不被尊重。”共同建立双方都能接受的社交边界。 | 建议在放松的环境下与伴侣进行“底线对话”。明确表达这种行为让你感到缺乏安全感，而不是一味指责。在中国感情观中，建立“边界感”是相互尊重的表现，你可以提议共同制定一些沟通准则。 | 找一个安静的时间坦诚沟通你的感受，设定明确的边界，解释这种行为如何影响了你们的信任，并听取他的看法。 | Focus on setting boundaries through open communication rather than "making" him stop, which can backfire. |
| 14 我们的邻居在其他邻居小孩可以听到的范围内使用不合适的语言。我该如何让他们注意到这种行为？ | 找一个轻松的时机，友好地提醒：“刚才孩子们在旁边玩，如果不小心学了那些词就尴尬了，能不能麻烦稍微留意一下？”通常这种提醒能让对方意识到影响。 | 最好的方式是利用“为了孩子”这个理由。你可以平和地对邻居说：“不好意思，家里小孩模仿能力强，希望能帮忙注意一下用词。”这种说法通常能给对方留面子，也更容易被接受。 | 可以先尝试礼貌地与邻居沟通，说明声音较大且有小孩在场；如果无效，可以联系居委会(RC)或社区调解中心(CMC)寻求帮助。 | Approach them calmly and privately, explaining that the kids are "parrots" who pick up new words easily. |